



redis

Need Some Cache?

Redis In Depth

Chris Meadows
Architect & Developer
Terenine Technology Solutions
Email: meadoch1@gmail.com
Twitter: @meadoch1

C H A T T A N O O G A - A U G 1 7 - 1 9 , 2 0 1 1

Agenda

What's Redis?

How do I make it work?

What's under the hood?

Q&A

Credit Where Due

Peter Cooper – <http://coder.io> or peterc (Twitter)

<http://redis.io> - one of the best OSS documentation sites

What is Redis?

“an...advanced key-value
store...” - from redis.io

(aka a NoSQL data store)

NoSQL?

A loose category of non-relational data storage technologies

Other examples

Memcached, MongoDB, RavenDB

CouchDB, Riak, Cassandra

Memcache

Used (largely) as a network based
hash/dictionary

Became popular largely because it is
simple to use and
fast

Step it up!

Add

- Configurable disk persistence options
- Multiple native data structure options
- Native support for replication
- Native support for pub/sub
- ...more

...and you get
Redis

Reasons to Consider Redis

- Functionally
 - Speed (more on this later)
 - Flexibility
 - Low hardware demand
 - Disk persistence guarantees (if configured)
 - Low development overhead
- “Politically”
 - Price (free – BSD licensing)
 - Actively developed
 - Good community (list, IRC & wiki)
 - Used broadly (GitHub, Twitter, StackOverflow...)
 - Solid Corporate backing by VMWare

And now for something completely
different...

Example Application Usage


Social Network Storage

Users

have names, can follow others, and be followed

Posts

are things like messages, photos, etc



Example Application: Foundation

```
1  using System;
2  using Sider; //https://github.com/chakrit/sider
3
4  namespace SimpleRedisExamples
5  {
6      public class BasicSocialNetwork
7      {
8
9          RedisClient Client;
10
11         public BasicSocialNetwork ()
12         {
13             Client = new RedisClient("localhost", 6379);
14
15             //CreateDemoData\(\);
16         }
17     }
```

Now back to your
regularly scheduled program

How do I make it work?

Data Structure Types

Type	Key	Value
String	login.html TotalLogins	"<html><head>..." 3949
List	user:1:messages	[1, 2, 3, 4]
Set	user:1:contacts	{100, 107, 304, 238}
Hash	user:1:settings	visibility -> contacts only name -> john
Sorted Set	userActivity	1 -> 20 2 -> 34 3 -> 102

Strings

Interaction from command line

client utility key value
↓ ↓ ↓
./redis-cli SET "mykey" 12345
 key
 ↓
./redis-cli GET "mykey" → 12345 returns value

Strings

GET

SET

MGET

MSET

SETNX

MSETNX

GETBIT

SETBIT

SETRANGE

GETRANGE

GETSET

SETEX

INCR

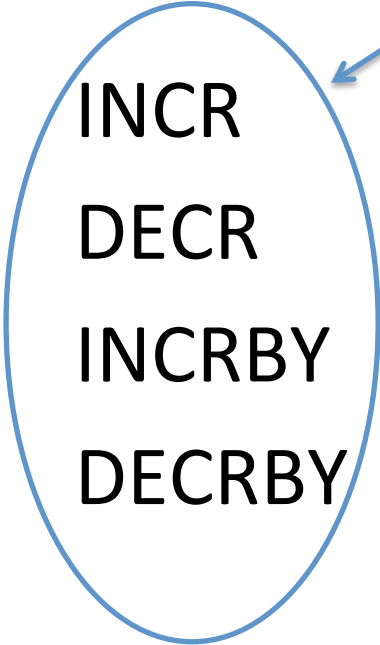
DECR

INCRBY

DECRBY

STRLEN

Work on
strings
which
represent
integers



Unique IDs

INCR next_post_id


If next_post_id doesn't exist or doesn't contain a number, it will be set to 0, incremented, and 1 will be returned.

returns → 1

INCR next_post_id

returns → 2

Example Application




user:1:name → joe

username:joe → 1

So we can do a two way
lookup

post:1:content → hello world

post:1:user → 1



SET user:1:name joe

SET username:joe 1


SET post:1:content "hello world"

SET post:1:user 1

Example App: Strings

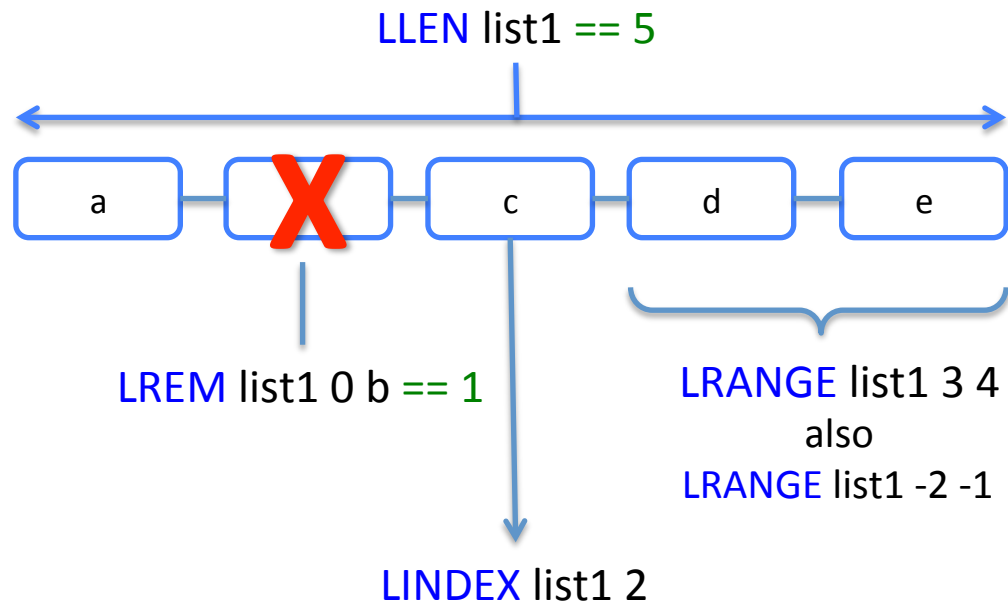
```
18 public long CreateNewUser(string userName)
19 {
20     long uid = Client.Incr("next_user_id");
21     Client.Set(string.Format("user:{0}:name", uid), userName);
22     Client.Set(string.Format("username:{0}", uid.ToString()), uid.ToString());
23     return uid;
24 }
```

Lists

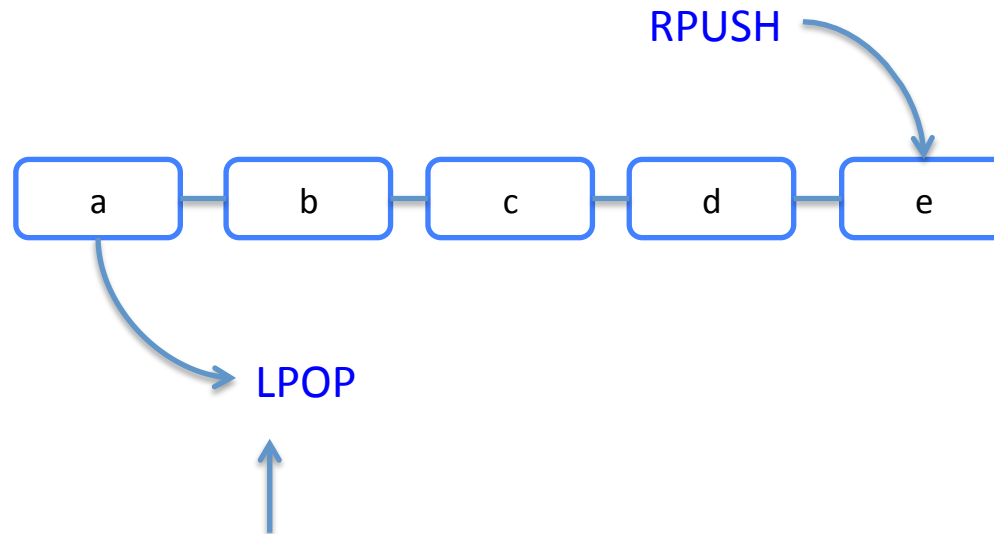


```
RPUSH list1 a -> 1
RPUSH list 1 b -> 2
LPOP list1 -> "a"
RPUSH list1 c -> 2
LPOP list1 -> "b"
LPOP list1 -> "c"
LPOP list1 -> (nil)
```

Lists




List Use Case Queue

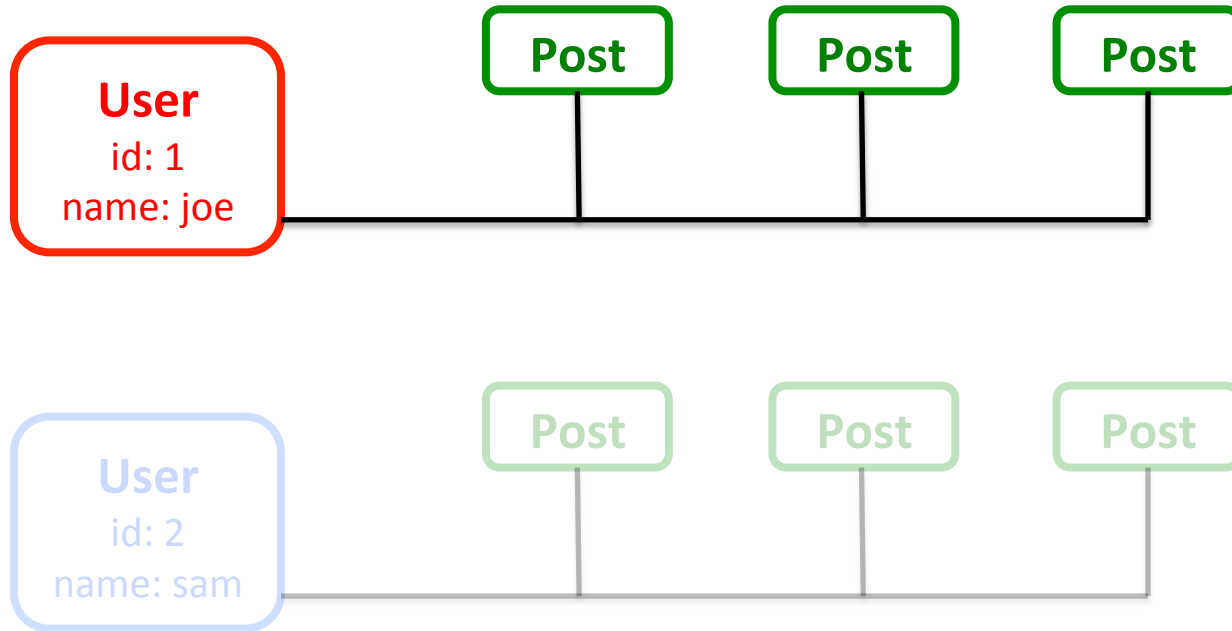


Or **BLPOP** to block
(wait) until something
can be popped

RPUSH the_q a
...
RPUSH the_q e
LPOP the_q
LPOP the_q
LPOP the_q



user:1:posts → [3, 2, 1]



user:1:posts → [3, 2, 1]

LPUSH user:1:posts 1

LPUSH user:1:posts 2

LPUSH user:1:posts 3

Example App: List

```
26 public long CreateNewPost(string postContent, long uid)
27 {
28     long pid = Client.Incr("next_post_id");
29     Client.Set(string.Format("post:{0}:content", pid), postContent);
30     Client.Set(string.Format("post:{0}:user", pid), uid.ToString());
31     Client.LPush(string.Format("user:{0}:posts", uid), pid.ToString());
32     Client.LPush("posts:global", pid.ToString());
33     return pid;
34 }
```

Sets

Holds an unordered group of items
accessed by a string key

contains:aba

abacus cabal baba hello teabag
base cabaret database

contains:ase

vase decease baseline case
database phase

Sets

Holds an unordered group of items
accessed by a string key

contains:aba

abacus cabal baba hello teabag
base cabaret database

contains:ase

vase decease baseline case
database phase

SADD contains:ase suitcase



Sets

Holds an unordered group of items
accessed by a string key

SREM contains:aba hello

contains:aba

abacus cabal baba ~~hello~~ teabag
base cabaret database

contains:ase

vase decease suitcase baseline
case database phase

Sets

Holds an unordered group of items
accessed by a string key

contains:aba

abacus cabal baba teabag **base**
cabaret database

SMOVE contains:aba contains:ase base

contains:ase

vase decease suitcase baseline
case database phase



Sets

contains:aba

abacus cabal baba teabag cabaret
database

SCARD contains:aba == 6

SISMEMBER contains:aba chips == 0 (meaning false)

SRANDMEMBER contains:aba == "teabag"

contains:ase

vase decease suitcase baseline
case database phase base

SMEMBERS contains:ase == vase, decease, suitcase,
baseline, case, database
phase, base

Sets

contains:aba

abacus cabal baba teabag cabaret

database

vase decease

suitcase baseline case phase base

contains:ase

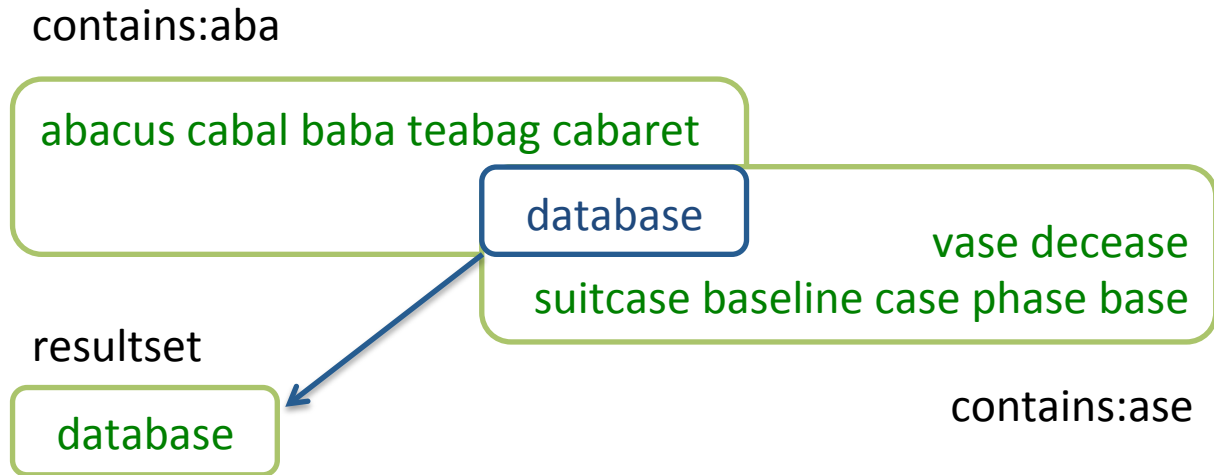
SINTER contains:aba contains:ase == database

This is a very simple example. **SINTER** can take any number of arguments.

UNION will join the sets together (not repeating duplicates).

SDIFF will return elements not in common between sets.

Sets



SINTERSTORE resultset contains:aba contains:ase

SUNIONSTORE does the same for set unions.


SDIFFSTORE does the same for set diffs.

Set Use Cases

Tags

Attributes

Property Bag




user:1:follows \longrightarrow {2, 3, 4}

SADD user:1:follows 2

SADD user:1:follows 3

SADD user:1:follows 4



user:1:followed_by \longrightarrow {2, 3}

SADD user:1:followed_by 2

SADD user:1:followed_by 3

Example App: Set

```
36 public void FollowUser(long followerId, long watchedId)
37 {
38     Client.SAdd(string.Format("user:{0}:follows", followerId),
39                 watchedId.ToString());
40     Client.SAdd(string.Format("user:{0}followed_by", watchedId),
41                 followerId.ToString());
42 }
```

Sorted Sets

Sorry – not enough time!

Work similarly to sets but each element has a “rank” or “score” associated with it to be used for sort order.

Hashes

product:1

Key	Value
created_at	102374657
product_id	1
name	Yo-yos
available	10

HSET product:1 created_at 102374657

HSET product:1 product_id 1

HSET product:1 name "Yo-yos"

HSET product:1 available 10

HGET product:1 name == Yo-yos

HLEN product:1 == 4

HKEYS product:1 == created_at, product_id,
name, available

HGETALL product:1 == created_at => 102374657
product_id => 1

...

And more...

HVALS HEXISTS HINCRBY HMGET HMSET

Hash Use Case

Session Storage

Session	8d3e4
created_at	102374657
user_id	1

Is essentially a hash

HSET session:8d3e4 created_at 102374657

HSET session:8d3e4 user_id 1


OR

HMSET session:8d3e4 created_at 102374657 session:8d3e4 user_id 1

For session time out let Redis automatically expire it in 24 hours

EXPIRE session:8d3e4 86400

Publish/Subscribe



Also PSUBSCRIBE supports wildcards

Enough Commands?

Redis also has commands for
Key Manipulation
Transactions
Connection Settings
Server Management

124 commands in all
(and counting)

What's Under the Hood?

Written in ANSI C (~20k lines)

Runs on most POSIX systems

Direct binary access for most languages

Event loop processing model

Installation

- Runs on POSIX systems
 - i.e. Linux, Unix, Solaris, OSX
 - Windows: Cygwin or fork of source
 - Not officially supported
- Options to get it running
 - Install from source (uses Make)
 - Install using package manager
 - Apt-get, homebrew and other packages available
- Once installed
 - Redis-server
 - Redis-cli

Configuration Options

General server settings – port, timeouts, log levels, etc

Snapshotting – when to save to disk, where

Replication – master address, password

Security – password (optional), command renaming

Limits – clients, memory

Append only mode – yes/no, fsync behavior

Slow log – parameters

Virtual Memory – yes/no, settings

Memory optimization settings – structure compression
parameters

Performance

Highly dependent on hardware, configuration,
and operation complexity

Also impacted by client library and language

For basic ops (GET, SET, LPUSH,
RPOP, HSET, HGET, etc)

Low end – 2000-5000 ops/sec

High end – 120,000+ ops/sec

Performance Example 1

On this machine

MacBook Pro i7 with 8Gb memory

Redis benchmark (3 byte packets)

MSET 36900.37 rps (99.89% <= 2ms)

SET 54644.81 rps (98.78% <= 1ms)

GET 56179.77 rps (99.24% <= 1ms)

INCR 54054.05 rps (98.47% <= 1ms)

LPUSH 55248.62 rps (99.93% <= 2ms)

LPOP 56497.18 rps (99.77% <= 1ms)

SADD 54945.05 rps (99.83% <= 2ms)

SPOP 55555.55 rps (100% <= 1ms)

Avg 53k
rps

Performance Example 2

On this machine

MacBook Pro i7 with 8Gb memory


Erlang client using Basho Bench

30 clients, 30000 keys

Simultaneous (random) k/v, hash, & list
access

½ single integer payload, ½ 2705 byte string
payload


Performance Example 2



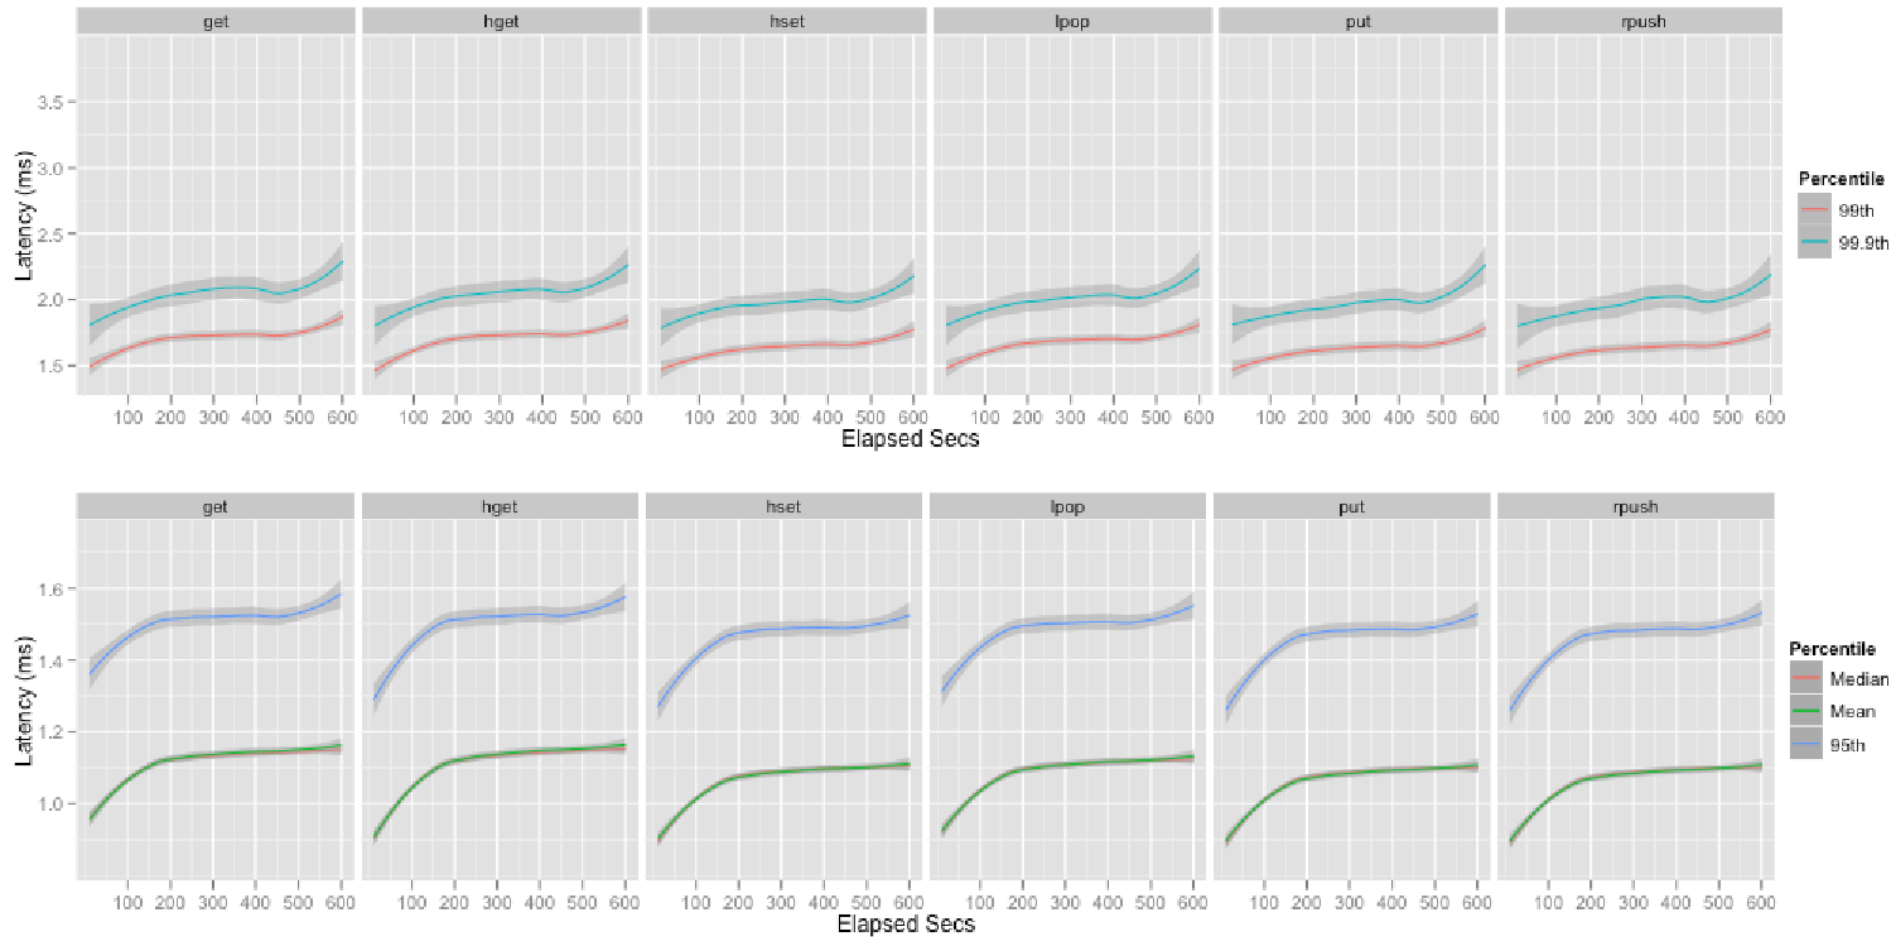
With snapshotting on

Average between 30,000-25,000 ops/sec across 10 minutes

With snapshotting off



Performance Example 2



Latencies average between <1-1.1ms
99.9% of all latencies < 2.0-2.3 ms

Languages

Most

(worth using and some that aren't)

C, C#, C++, Clojure, Common Lisp, Erlang, Go, Haskell, Io, Java, Lua, Node.js, Objective-C, Perl, PHP, Python, Ruby, Scala, and more...

<http://redis.io/clients>

More info

We just scratched the surface

Dig into redis.io for more depth

Google “redis”

I’d be glad to talk

Chris Meadows

meadoch1@gmail.com

@meadoch1 (Twitter)

<https://github.com/meadoch1/RedisIntro>